

## Assignment- 2

### Transport Layer and Network Simulations using NS-3

#### PART 3

### Understanding TCP Congestion Window using NS-3

**1. For each of the given congestion control algorithms, perform the simulations and answer the following questions.**

- a. Newreno**
- b. Highspeed**
- c. Veno**
- d. Vegas**

**Q1. Plot the cwnd vs. time graph, and describe what you observed, like slow start and congestion avoidance, in detail.**

Using Gnuplot with the given script, plotting the graph between cwnd vs. time for each algorithm.

**Steps to generate demo.cwnd file:-**

- Copy Demo.cc file in **scratch/** folder.

```
> cd scratch/
> ls
CMakeLists.txt Demo.cc first.cc nested-subdir q2.cc scratch-simulator.cc subdir
~/Downloads/ns-allinone-3.40/ns-3.40/scratch >
```

- Run **.ns3 run scratch/Demo.cc**

```
> ./ns3 run scratch/Demo.cc
[0/2] Re-checking globbed directories...
[1/1] Linking CXX executable /home/mafia/Downloads/ns-allinone-3.40/ns-3.40/build/scratch/ns3.40-Demo-debug
~/Downloads/ns-allinone-3.40/ns-3.40 >
```

By changing the Demo.cc file in this line :

“Config::SetDefault("ns3::TcpL4Protocol::SocketType", TypeIdValue (<**Algorithm Name**>::GetTypeId()));”

We can observe for different algorithms.

- After that, we will get an executable file **ns3.40-Demo-debug** at **build/scratch/<file>**.

```
> cd Downloads/ns-allinone-3.40/ns-3.40/
> ls
__pycache__ Basant-2-1.pcap cmake-cache demo-file-1-1.pcap examples scratch utils
AUTHORS Basant-3-0.pcap CMakeLists.txt demo-file-2-0.pcap LICENSE setup.cfg utils.py
Basant-0-0.pcap bindings contrib demo-file-2-1.pcap ns3 setup.py VERSION
Basant-1-0.pcap build CONTRIBUTING.md demo-file-3-0.pcap pyproject.toml src
Basant-1-1.pcap build-support demo-file-0-0.pcap demo.cwnd README.md test.py
Basant-2-0.pcap CHANGES.md demo-file-1-0.pcap doc RELEASE_NOTES.md testpy-output
> cd build
> ls
examples include lib scratch src utils
> cd scratch
> ls
nested-subdir ns3.40-Demo-debug subdir
~/Downloads/ns-allinone-3.40/ns-3.40/build/scratch >
```

Running this file will provide three files as follows: one is demo.cwnd:

```
> ls
examples include lib scratch src utils
> cd scratch
> ls
demo-file-0-0.pcap demo.cwnd ns3.40-Demo-debug
demo-file-1-0.pcap nested-subdir subdir
~/Downloads/ns-allinone-3.40/ns-3.40/build/scratch >
```


>> graph.pg -> uses gnuplot to generate graph.

```
3 set term postscript eps color
2 set output 'cwnd.eps'
1 set ylabel 'cwnd'
4 set xlabel 'time'
1 plot 'demo.cwnd'
```

Then run: **>> gnuplot graph.gp**


Doing the above steps for all the algorithms.

#### a. Tcp NewReno



The NewReno algorithm uses a fast retransmit technique to recover from loss events quickly. It immediately retransmits when a loss event is detected and starts a timer. Suppose the timer expires before an ACK is received for the retransmitted packet. In that case, NewReno assumes that the packet has been lost again and retransmits it repeatedly until an ACK is received for the retransmitted packet.


## b. Tcp HighSpeed



In the Highspeed algorithm, a slow start restart is used to recover from the loss events and avoid congestion quickly. When Highspeed detects a loss event, it retransmits the lost packet and resets the cwnd to half its previous value.


Highspeed then enters a slow start phase, doubling the cwnd each time it receives an ACK. This continues until the cwnd reaches its previous value.

### c. Tcp Veno



In the Tcp Veno algorithm, a combination of slow start and congestion avoidance is used. In the slow start phase, it doubles the cwnd each time it receives an ACK. In the congestion avoidance mode, it increases the cwnd by one each time it receives ACK. But this algorithm also uses an Adaptive exponential increase/multiplicative decrease (AIMD) to adjust the rate at which the cwnd will be increased.

#### d. Tcp Vegas



The Tcp Vegas algorithm uses rate-based pacing to avoid congestion. It maintains a running estimate of the available bandwidth and uses this estimate to adjust the rate at which it sends packets. If it detects that the network is congested, it will slow down the rate at which it sends packets.

**Q2. Find the average throughput for each of the congestion control algorithms using tshark from the pcap files generated, and state which algorithm performed the best.**

**Command:** tshark -r demo-file-1-0.pcap -qz io,stat,1,BYTES

**a. Tcp NewReno**

```
=====
| IO Statistics |
| Duration: 18.999916 secs |
| Interval: 1 secs |
| Col 1: BYTES |
|-----|
| Interval | 1 | BYTES |
|-----|
| 0 <> 1 | 634884 |
| 1 <> 2 | 656156 |
| 2 <> 3 | 579056 |
| 3 <> 4 | 621386 |
| 4 <> 5 | 622864 |
| 5 <> 6 | 652898 |
| 6 <> 7 | 655580 |
| 7 <> 8 | 587708 |
| 8 <> 9 | 623674 |
| 9 <> 10 | 631268 |
| 10 <> 11 | 655066 |
| 11 <> 12 | 606032 |
| 12 <> 13 | 581242 |
| 13 <> 14 | 614302 |
| 14 <> 15 | 589370 |
| 15 <> 16 | 631798 |
| 16 <> 17 | 627064 |
| 17 <> 18 | 654294 |
| 18 <> Dur | 632360 |
|-----|
=====
```

**Average Throughput:** Total Bytes/ Duration

$$= \frac{11857002}{18.999916}$$

$$= \mathbf{624055.4958 \text{ Bytes/sec}}$$

## b. Tcp Highspeed

=====			
IO Statistics			
Duration: 18.999132 secs			
Interval: 1 secs			
Col 1: BYTES			
-----			
		1	
Interval		BYTES	
-----			
0 <>	1	653810	
1 <>	2	659430	
2 <>	3	656704	
3 <>	4	658756	
4 <>	5	655390	
5 <>	6	657086	
6 <>	7	661036	
7 <>	8	66804	
8 <>	9	606828	
9 <>	10	657716	
10 <>	11	654760	
11 <>	12	658556	
12 <>	13	653886	
13 <>	14	656788	
14 <>	15	656096	
15 <>	16	657664	
16 <>	17	657568	
17 <>	18	656104	
18 <>	Dur	658346	
=====			

**Average Throughput: Total Bytes/ Duration**

$$= \frac{11843328}{18.999132}$$

$$= \mathbf{623361.5304 \text{ Bytes/sec}}$$



### c. Tcp Veno

=====		
IO Statistics		
Duration: 18.999724 secs		
Interval: 1 secs		
Col 1: BYTES		
-----		
	1	
Interval	BYTES	
-----		
0 <> 1	634240	
1 <> 2	656156	
2 <> 3	579056	
3 <> 4	618972	
4 <> 5	622810	
5 <> 6	652952	
6 <> 7	657140	
7 <> 8	581794	
8 <> 9	623022	
9 <> 10	631260	
10 <> 11	656494	
11 <> 12	606498	
12 <> 13	442436	
13 <> 14	644692	
14 <> 15	572312	
15 <> 16	614150	
16 <> 17	626300	
17 <> 18	654966	
18 <> Dur	628028	
-----		
=====		

**Average Throughput:** Total Bytes/ Duration

$$= \frac{11703278}{18.999724}$$

$$= \mathbf{615970.9478 \text{ Bytes/sec}}$$

#### d. Tcp Vegas

```
=====
```

IO Statistics		
Duration: 18.998761 secs		
Interval: 1 secs		
Col 1: BYTES		
-----		
	1	
Interval	BYTES	
-----		
0 <> 1	223458	
1 <> 2	421666	
2 <> 3	634590	
3 <> 4	614282	
4 <> 5	627404	
5 <> 6	636072	
6 <> 7	645298	
7 <> 8	632448	
8 <> 9	625254	
9 <> 10	628444	
10 <> 11	634620	
11 <> 12	641876	
12 <> 13	619760	
13 <> 14	635256	
14 <> 15	613250	
15 <> 16	632564	
16 <> 17	618744	
17 <> 18	639434	
18 <> Dur	641232	
	.....	

```
=====
```

**Average Throughput:** Total Bytes/ Duration

$$= \frac{11365652}{18.998761}$$

$$= \mathbf{598231.2215 \text{ Bytes/sec}}$$

Tcp NewReno Performed best among all the algorithms with an average throughput of **624055.4958** Bytes/sec.

### Q3. How many times did the TCP algo reduce the cwnd and why?

Counting the approx. number of peaks for each algorithm.

- a. Newreno:- 81
- b. Highspeed:- 75
- c. Veno:- 79
- d. Vegas:- 67

There can be various reasons for reducing cwnd. Some are mentioned below.


- **Packet loss:** TCp reduces cwnd when it detects packet loss.
- **Explicit Congestion Notification (ECN):** Enabled ECN can lead to cwnd reduction when network congestion is signed without packet loss.
- **Timeout:** A timeout event also indicates that congestion leads to cwnd reduction.
- **Triple Duplicate ACKs:** After detecting three ACK, some TCP variants reduce cwd as a congestion signal.

**Q4. Check the effect of changing the bandwidth and latency of point-to-point connection and explain its effect on average throughput.**

**a. Newreno**

- Increased bandwidth typically increases throughput, while decreased bandwidth produces lower throughput.
- Latency: As Newreno adjusts to network conditions, higher latency can result in lower throughput, and lower latency can result in higher throughput.

Simulation:- bw = 10mbps, latency = 1ms



```
=====
```

IO Statistics		
Duration: 18.998698 secs		
Interval: 1 secs		
Col 1: BYTES		
-----		
	1	
Interval		BYTES
-----		
0 <>	1	1283478
1 <>	2	1191646
2 <>	3	1271578
3 <>	4	1236798
4 <>	5	1249182
5 <>	6	1255004
6 <>	7	1188464
7 <>	8	1213552
8 <>	9	1275466
9 <>	10	1247658
10 <>	11	1244948
11 <>	12	1265906
12 <>	13	1218598
13 <>	14	34338
14 <>	15	997378
15 <>	16	1238250
16 <>	17	1240786
17 <>	18	1251930
18 <>	Dur	1274142
.....		

```
=====
```

**Average Throughput:** Total Bytes/ Duration

$$= \frac{22179102}{18.998698}$$


$$= 1167401.156 \text{ Bytes/sec}$$

Average throughput increases on an increase in bandwidth and a decrease in the latency from the Q2 TcpNewReno throughput (624055.4958).

## b. HighSpeed

- Throughput should increase when bandwidth is increased, but there may be diminishing returns at very high bandwidth levels. Throughput will decline as bandwidth is reduced.
- Low-latency environments should allow high speed to remain steady and high throughput. Because round-trip times increase with increased latency, throughput may suffer.

Simulation:- bw = 10mbps, latency = 5ms



IO Statistics		
Duration: 18.994624 secs		
Interval: 1 secs		
Col 1: BYTES		
-----		
	1	
Interval	BYTES	
-----		
0 <> 1	1057656	
1 <> 2	923076	
2 <> 3	827538	
3 <> 4	1219670	
4 <> 5	614472	
5 <> 6	250792	
6 <> 7	1000368	
7 <> 8	1125776	
8 <> 9	838398	
9 <> 10	901156	
10 <> 11	901894	
11 <> 12	1104676	
12 <> 13	1044064	
13 <> 14	1022382	
14 <> 15	888132	
15 <> 16	1065078	
16 <> 17	496590	
17 <> 18	371878	
18 <> Dur	847452	
.....		

**Average Throughput:** Total Bytes/ Duration

$$= \frac{16501048}{18.994624}$$


$$= \mathbf{868722.0131 \text{ Bytes/sec}}$$

Avg. throughput increase on the increase in bandwidth from 5bytes to 10 bytes  
(previous Throughput=623361.5304 Bytes/sec)

### c. Veno

- Throughput is typically improved when bandwidth is increased, while throughput is typically decreased when bandwidth is decreased.
- Throughput may be less susceptible to changes in latency since Veno can manage varying latencies, although throughput can still be negatively impacted by very high latency.

Simulation:- bw = 10mbps, latency = 1ms





```
=====
```

IO Statistics		
Duration: 18.981829 secs		
Interval: 1 secs		
Col 1: BYTES		
-----		
	1	
Interval	BYTES	
-----		
0 <> 1	306654	
1 <> 2	96222	
2 <> 3	416116	
3 <> 4	471444	
4 <> 5	374256	
5 <> 6	248084	
6 <> 7	305758	
7 <> 8	312252	
8 <> 9	395240	
9 <> 10	486170	
10 <> 11	678914	
11 <> 12	433428	
12 <> 13	252296	
13 <> 14	410952	
14 <> 15	386834	
15 <> 16	363266	
16 <> 17	429682	
17 <> 18	547078	
18 <> Dur	314028	
	....	

```
=====
```

**Average Throughput:** Total Bytes/ Duration

$$= \frac{7228674}{18.981829}$$

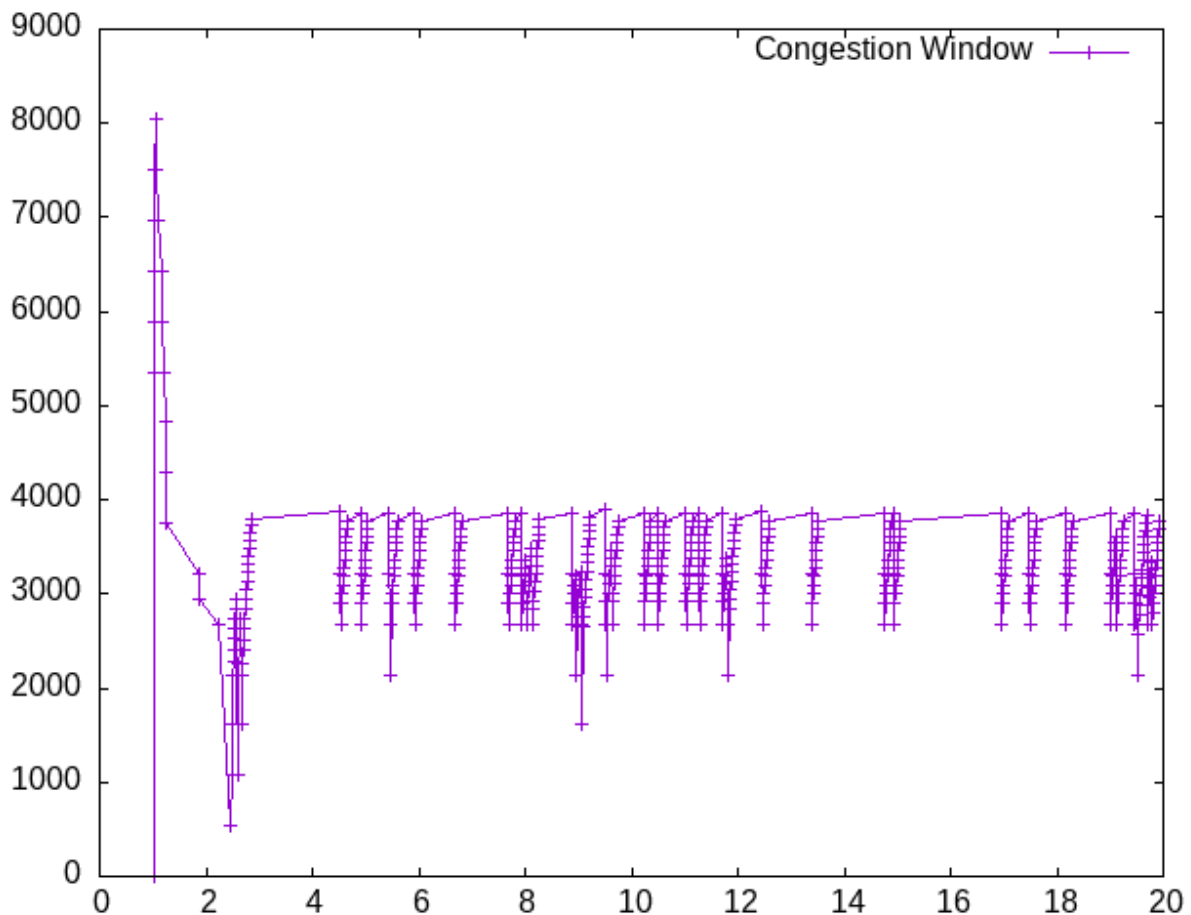
$$= \mathbf{380820.7312 \text{ Bytes/sec}}$$

Average throughput decreases on an increase in latency from 2ms to 10ms.  
(prev throughput =615970.9478 Bytes/sec).

#### d. Vegas

- Increased bandwidth leads to increased throughput, whereas decreased bandwidth may temporarily lower performance before Vegas adapts.
- Vegas is built to accommodate changes in bandwidth and latency, thus, changes in latency have less of an impact on throughput, although highly high latency may cause some throughput drop as Vegas adapts.

Simulation:- bw = 10mbps, latency = 1ms



=====			
IO Statistics			
Duration: 18.992449 secs			
Interval: 1 secs			
Col 1: BYTES			
-----			
		1	
Interval		BYTES	
-----			
0 <>	1	221688	
1 <>	2	105388	
2 <>	3	261608	
3 <>	4	252296	
4 <>	5	250084	
5 <>	6	256308	
6 <>	7	251714	
7 <>	8	242246	
8 <>	9	237932	
9 <>	10	252350	
10 <>	11	241408	
11 <>	12	256308	
12 <>	13	256362	
13 <>	14	252940	
14 <>	15	260964	
15 <>	16	258186	
16 <>	17	254484	
17 <>	18	256642	
18 <>	Dur	228370	
-----			
=====			

**Average Throughput:** Total Bytes/ Duration

$$= \frac{4597278}{18.992449}$$

$$= \mathbf{242058.1990 \text{ Bytes/sec}}$$

Average throughput decreases drastically on the decrease in bandwidth and an increase in latency. (prev throughput = 598231.2215 Bytes/sec).

#### **Q5. Explain in short what is the effect of changing the default MTU size.**

**MTU:-** The maximum transmission unit is the largest size frame or packet in bytes that can be transmitted across a data link. It is most used for packet size on an Ethernet network using the Internet Protocol.

To achieve the same network speed, the operating system can deliver fewer packets of a bigger size when using large MTU sizes. If the workload permits sending large messages, the larger packets significantly reduce the processing the operating system needs. The greater MTU size won't help if all of the workload consists of sending brief messages.

Effects of changing the default MTU size:-

- Affects network performance and compatibility.
- Larger MTUs can improve performance but may lead to fragmentation.

2. Using the Demo.cc file as reference, create a new scenario topology consisting of 4 nodes as shown below.


Q1. Find the average throughput for the given parameters.

bw1 = 5mbps, latency1 = 2ms

bw2 = 5mbps, latency2 = 2ms

bw3 = 5mbps, latency3 = 2ms

Using TcpNewReno algorithm for performing this simulation.



IO Statistics		
Duration: 18.992858 secs		
Interval: 1 secs		
Col 1: BYTES		
-----		
	1	
Interval	BYTES	
-----		
0 <> 1	344574	
1 <> 2	551966	
2 <> 3	413442	
3 <> 4	341736	
4 <> 5	295480	
5 <> 6	290026	
6 <> 7	293672	
7 <> 8	414774	
8 <> 9	491942	
9 <> 10	597128	
10 <> 11	426882	
11 <> 12	227648	
12 <> 13	434202	
13 <> 14	304194	
14 <> 15	453760	
15 <> 16	453718	
16 <> 17	561746	
17 <> 18	376884	
18 <> Dur	252438	
.....		

**Avg. throughput = Total Bytes/ Duration**

$$= \frac{7526212}{18.992858}$$

$$= \frac{512338}{3}$$

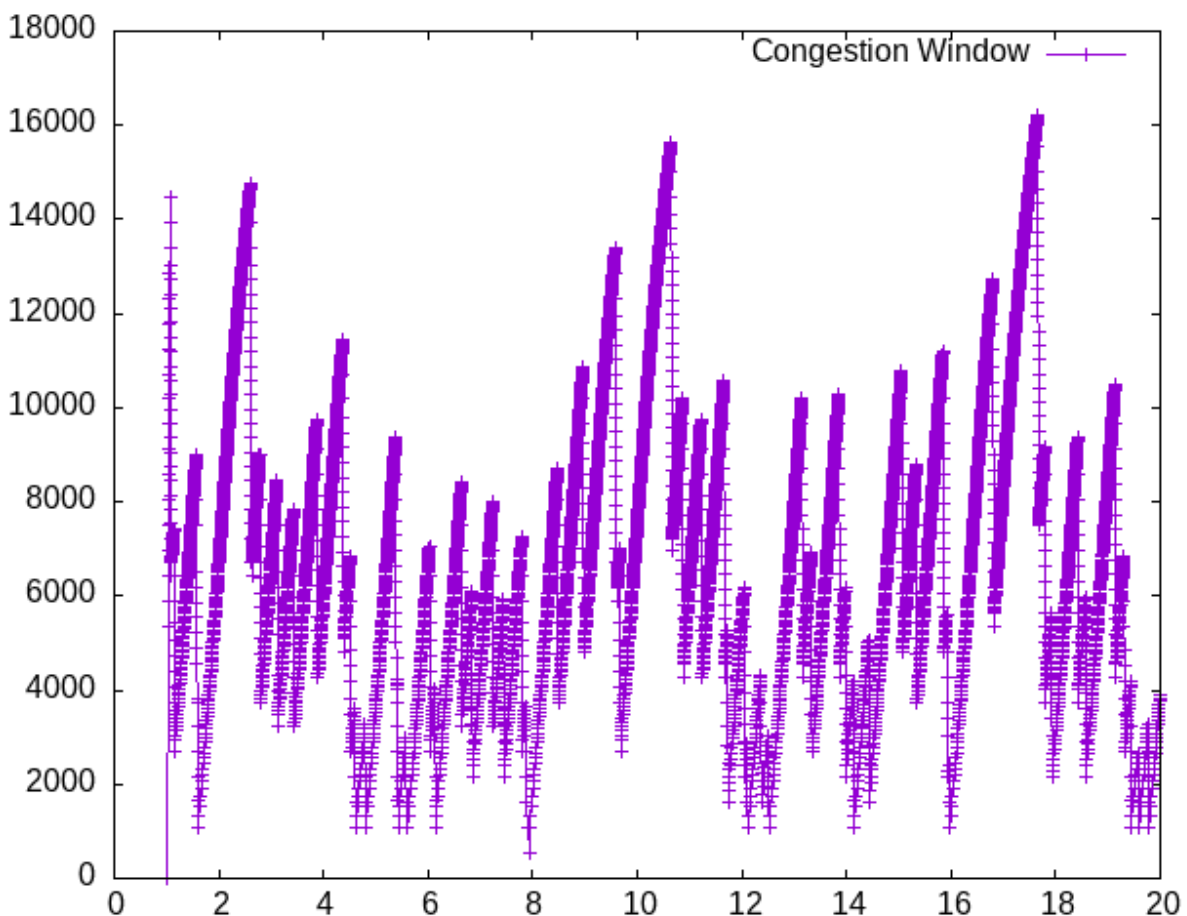
$$= 396265.3751 \text{ bps}$$

**Q2. Now change the values of latency2 from 2ms to 10ms and plot and compare the cwnd shape. Do this for two different congestion control algorithms.**

For this part, I used the modified version of [Demo.cc](#) script to find the throughput. The screenshot shows the throughput output for delay values at latency2 2ms,4ms,6ms,8ms, and 10ms, respectively, for both TcpNewReno and TcpVegas.

**a. TcpNewReno**

**1. Latency2=2ms**




**Average Throughput:** Total Bytes/ Duration

$$= \frac{7526212}{18.992858}$$

$$= 396265.3751 \text{ bps}$$

## 2. 4ms




**Average Throughput: Total Bytes/ Duration**

$$= \frac{5993720}{18.990784}$$

$$= 315612.0358 \text{ bps}$$



### 3. 6ms

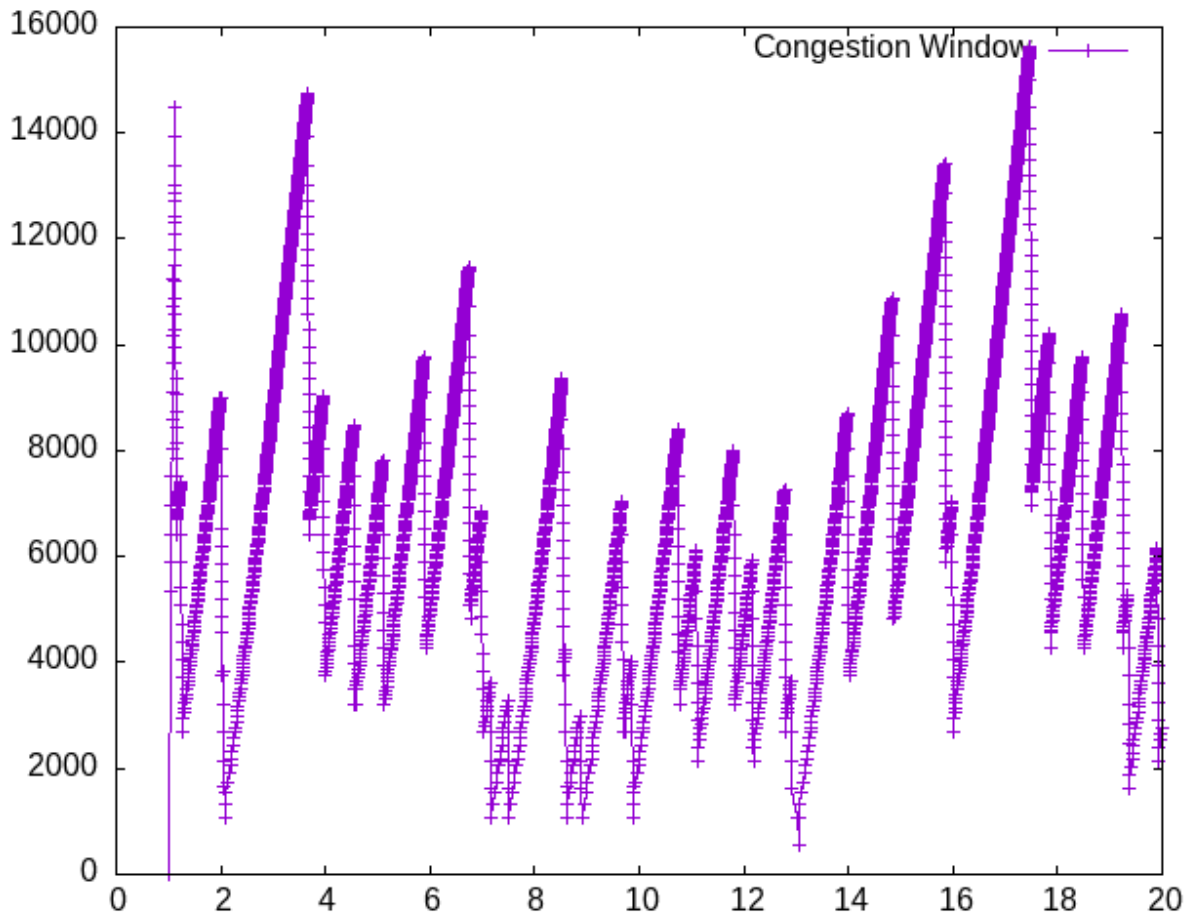


**Average Throughput:** Total Bytes/ Duration

$$= \frac{5090324}{18.981917}$$

$$= 268167.0139 \text{ bps}$$

#### 4. 8ms

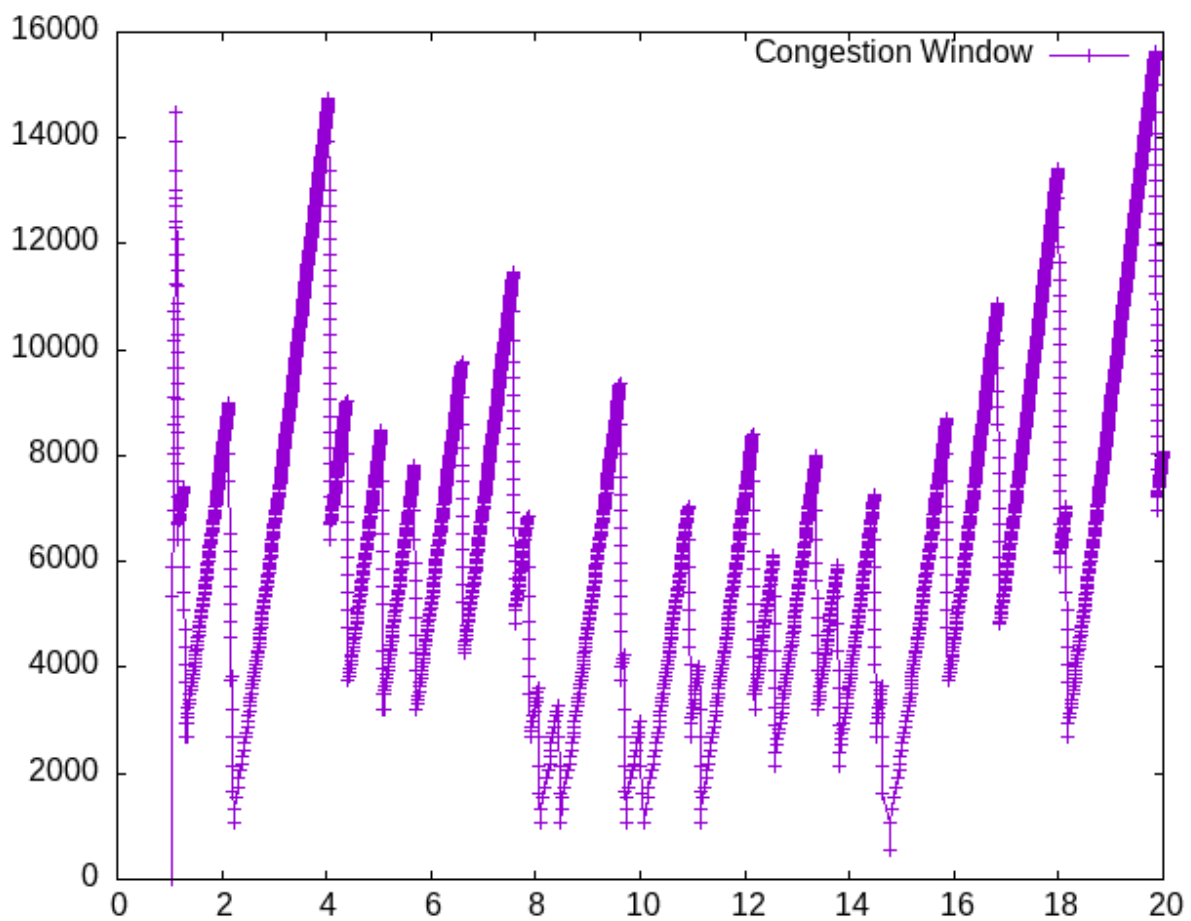


**Average Throughput:** Total Bytes/ Duration

$$= \frac{4484968}{18.963590}$$

$$= 236504.1640 \text{ bps}$$

## 5. 10ms




**Average Throughput:** Total Bytes/ Duration

$$= \frac{3895242}{18.985334}$$

$$= 205171.1073 \text{ bps}$$

The graph uses the values from 2ms to 10ms of latency<sup>2</sup> using a Python script.


```
1 import matplotlib.pyplot as plt
2
3 x = [2,4,6,8,10]
4 y = [396265.3751,315612.0358,268167.0139,236504.1640,205171.1073]
5
6 plt.plot(x,y)
7
8 plt.xlabel('Delay(ms)')
9 plt.ylabel('Avg. Throughput (bps)')
10
11 plt.title('Avg. Throughput changing with delay')
12
13 plt.show()
```



## B. TcpVegas

Repeating the same above step for this Algorithm.

### 1. 2ms




**Average Throughput:** Total Bytes/ Duration

$$= \frac{4862082}{18.988787}$$

$$= 256050.1627 \text{ bps}$$

## 2. 4ms




**Average Throughput:** Total Bytes/ Duration

$$= \frac{1364656}{3.997699}$$

$$= 341360.3676 \text{ bps}$$

### 3. 6ms




**Average Throughput:** Total Bytes/ Duration

$$= \frac{3029296}{18.977834}$$

$$= 159622.8526 \text{ bps}$$

#### 4. 8ms




**Average Throughput:** Total Bytes/ Duration

$$= \frac{2623984}{18.987187}$$

$$= 138197.6172 \text{ bps}$$



## 5. 10ms



**Average Throughput:** Total Bytes/ Duration

$$= \frac{1826194}{8.334672}$$

$$= 219108.0825 \text{ bps}$$

The graph is made using the values from 2ms to 10ms of latency<sup>2</sup> using the same Python script.

